

graph_analysis

January 22, 2026

1 Graph Analysis dengan PySpark GraphFrames

Nama: Muhammad Rizky Hajar

NIM: 24.55.2714

Prodi: PJJ Informatika

Konsentrasi: Big Data dan Predictive Analytics (BDPA)

Mata Kuliah: IT Platform and Infrastructure

1.1 Tugas 4: Graph Analysis - Flight Network

Konteks Tugas: - Menggunakan PySpark GraphFrames untuk analisis graph - Dataset: Flights dan Airports data - Analisis network penerbangan dan delay patterns - Visualisasi dengan Bokeh, NetworkX, dan Neo4j

1.1.1 Notebook Contents:

1. Introduction to GraphFrames
2. Setup & Import Libraries
3. Data Capture
4. Graph Construction
5. Graph Analysis & Queries
6. Graph Algorithms
7. Visualization
8. Conclusion

1.2 1. Introduction to GraphFrames

1.2.1 Apa itu GraphFrames?

GraphFrames adalah package Python untuk graph processing di Apache Spark. GraphFrames menyediakan API tingkat tinggi untuk graph analysis dan memanfaatkan kekuatan Spark DataFrames.

1.2.2 Komponen Graph:

- **Vertices (Nodes)**: Entitas dalam graph (dalam kasus ini: Airports)
- **Edges**: Koneksi antar vertices (dalam kasus ini: Flight routes)

1.2.3 Use Case:

Analisis network penerbangan untuk memahami:
- Bandara tersibuk - Pola delay penerbangan
- Hub airports dan connectivity - Optimal routes dan shortest paths

1.3 2. Setup & Import Libraries

```
[1]: # Import libraries
from pyspark.sql import SparkSession
from pyspark.sql.types import *
# Import specific functions to avoid conflicts with Python built-ins
from pyspark.sql.functions import (
    col, avg, count, sqrt, sum as spark_sum, lit, desc, row_number,
    when, length, concat, desc, asc, size, isnan
)
import pandas as pd
import numpy as np

# Visualization libraries
import matplotlib.pyplot as plt
import seaborn as sns
import networkx as nx
from bokeh.io import output_file, show
from bokeh.models import (
    GMapPlot, GMapOptions, ColumnDataSource, Circle, Line,
    Range1d, PanTool, WheelZoomTool, BoxSelectTool
)
import folium

# Neo4j connection (optional)
try:
    from py2neo import Graph, Node, Relationship
    NEO4J_AVAILABLE = True
except:
    NEO4J_AVAILABLE = False
    print("Neo4j library not available, skipping Neo4j visualizations")

# Warnings
import warnings
warnings.filterwarnings('ignore')

# Set matplotlib style
plt.style.use('seaborn-v0_8-darkgrid')
```

```

sns.set_palette("husl")

print("Libraries imported successfully!")

```

Libraries imported successfully!

```

[2]: # Initialize SparkSession with GraphFrames support
# GraphFrames 0.10.0 supports Spark 4.x!
spark = SparkSession.builder \
    .appName("FlightNetworkGraphAnalysis") \
    .config("spark.driver.memory", "4g") \
    .config("spark.sql.shuffle.partitions", "8") \
    .config("spark.jars.packages", "io.graphframes:graphframes-spark4_2.13:0.10.
˓→0") \
    .getOrCreate()

# Set log level
spark.sparkContext.setLogLevel("WARN")

print(f"Spark Version: {spark.version}")
print(f"Spark Session: {spark}")

# Import GraphFrames
from graphframes import GraphFrame
print(" GraphFrames 0.10.0 imported successfully!")

```

```

WARNING: Using incubator modules: jdk.incubator.vector
:: loading settings :: url = jar:file:/Users/riskihajar/projects/kuliah/dl-
tugas-2/venv/lib/python3.12/site-packages/pyspark/jars/ivy-
2.5.3.jar!/org/apache/ivy/core/settings/ivysettings.xml
Ivy Default Cache set to: /Users/riskihajar/.ivy2.5.2/cache
The jars for the packages stored in: /Users/riskihajar/.ivy2.5.2/jars
io.graphframes#graphframes-spark4_2.13 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-
parent-1c66a14f-15c5-4720-a1e1-6843dceeb4c;1.0
    confs: [default]
        found io.graphframes#graphframes-spark4_2.13;0.10.0 in central
        found io.graphframes#graphframes-graphx-spark4_2.13;0.10.0 in central
:: resolution report :: resolve 54ms :: artifacts dl 2ms
    :: modules in use:
        io.graphframes#graphframes-graphx-spark4_2.13;0.10.0 from central in
[default]
    io.graphframes#graphframes-spark4_2.13;0.10.0 from central in [default]
-----
|           |           modules           ||   artifacts   |
|       conf    | number| search|dwnlded|evicted|| number|dwnlded|
-----
|     default  |   2   |   0   |   0   |   0   ||   2   |   0   |

```

```

:: retrieving :: org.apache.spark#spark-submit-
parent-1c66a14f-15c5-4720-a1e1-6843dceeb4c
    confs: [default]
    0 artifacts copied, 2 already retrieved (0kB/2ms)
26/01/22 21:26:24 WARN NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).

Spark Version: 4.1.1
Spark Session: <pyspark.sql.session.SparkSession object at 0x130d57da0>
  GraphFrames 0.10.0 imported successfully!

```

1.4 3. Data Capture

1.4.1 Dataset Description:

1. **airports.dat** - OpenFlights airport database
 - Contains: Airport ID, Name, City, Country, IATA/ICAO codes, Latitude, Longitude, etc.
2. **flights.csv** - Flight performance data from Bureau of Transportation Statistics
 - Contains: Flight schedule, delays, origin/destination, dates, etc.

```
[3]: # Define file paths
tripdelaysFilePath = "../data/flights/flights.csv"
airportsnaFilePath = "../data/flights/airports.dat"

print("Data file paths defined")
```

Data file paths defined

1.4.2 3.1 Load Airports Data

```
[4]: # Define schema for airports data
# Format: AirportID, Name, City, Country, IATA, ICAO, Latitude, Longitude,
#          Altitude, Timezone, DST, Tz_database_time_zone, Type, Source
schemaString = "AirportID Name City Country IATA ICAO Latitude Longitude\u
    ↴Altitude Timezone DST Tz_database_time_zone Type Source"
fields = [StructField(field_name, StringType(), True) for field_name in\u
    ↴schemaString.split()]
airport_schema = StructType(fields)

# Load airports dataset
airportsna = spark.read.csv(
    airportsnaFilePath,
    header=False,
```

```

    schema=airport_schema,
    sep=',',
    quote='''',
    escape=''
)

# Create temp view for SQL queries
airportsna.createOrReplaceTempView("airports_na")

print(f"Airports loaded: {airportsna.count()} records")
airportsna.show(5)

```

Airports loaded: 7698 records

AirportID	Name	City	Country	IATA	ICAO	Type
Latitude	Longitude	Altitude	Timezone	DST	Tz_database_time_zone	
Source						
1	Goroka Airport	Goroka	Papua New Guinea			
GKA AYGA -6.081689834590001	145.391998291	5282	10	U		
Pacific/Port_Moresby	airport	OurAirports				
2	Madang Airport	Madang	Papua New Guinea	MAG	AYMD	
-5.20707988739	145.789001465	20	10	U		
Pacific/Port_Moresby	airport	OurAirports				
3	Mount Hagen Kagam...	Mount Hagen	Papua New Guinea			
HGU AYMH -5.826789855957031	144.29600524902344	5388	10	U		
Pacific/Port_Moresby	airport	OurAirports				
4	Nadzab Airport	Nadzab	Papua New Guinea	LAE	AYNZ	
-6.569803	146.725977	239	10	U		
Pacific/Port_Moresby	airport	OurAirports				
5	Port Moresby Jack...	Port Moresby	Papua New Guinea			
POM AYPY -9.443380355834961	147.22000122070312	146	10	U		
Pacific/Port_Moresby	airport	OurAirports				

only showing top 5 rows

[5]: # Filter USA airports and convert to Pandas for visualization

```

airportna_df = airportsna.toPandas()
airportna_usa_df = airportna_df[airportna_df['Country'] == 'United States']
airportna_usa_df[['Latitude', 'Longitude']] =_
    ↪airportna_usa_df[['Latitude', 'Longitude']].apply(pd.to_numeric)

```

```

print(f"USA Airports: {len(airportna_usa_df)} records")
print(f"\nSample USA airports:")
airportna_usa_df[['IATA', 'Name', 'City', 'Latitude', 'Longitude']].head()

```

USA Airports: 1512 records

Sample USA airports:

	IATA	Name	City	Latitude	Longitude
3212	BTI	Barter Island LRRS Airport	Barter Island	70.134003	-143.582001
3213	\N	Wainwright Air Station	Fort Wainwright	70.613403	-159.860001
3214	LUR	Cape Lisburne LRRS Airport	Cape Lisburne	68.875099	-166.110001
3215	PIZ	Point Lay LRRS Airport	Point Lay	69.732903	-163.005005
3216	ITO	Hilo International Airport	Hilo	19.721399	-155.048004

1.4.3 3.2 Load Flights Data

```

[6]: # Load departure delays data
departureDelays_raw = spark.read.csv(
    tripdelaysFilePath,
    header=True,
    inferSchema=True
)

departureDelays_raw.createOrReplaceTempView("departureDelays_raw")

print(f"Flights loaded: {departureDelays_raw.count()} records")
print(f"\nSchema:")
departureDelays_raw.printSchema()

```

Flights loaded: 274964 records

Schema:

```

root
|-- YEAR: integer (nullable = true)
|-- MONTH: integer (nullable = true)
|-- DAY: integer (nullable = true)
|-- DAY_OF_WEEK: integer (nullable = true)
|-- AIRLINE: string (nullable = true)
|-- FLIGHT_NUMBER: integer (nullable = true)
|-- TAIL_NUMBER: string (nullable = true)
|-- ORIGIN_AIRPORT: string (nullable = true)
|-- DESTINATION_AIRPORT: string (nullable = true)
|-- SCHEDULED_DEPARTURE: integer (nullable = true)
|-- DEPARTURE_TIME: integer (nullable = true)
|-- DEPARTURE_DELAY: integer (nullable = true)
|-- TAXI_OUT: integer (nullable = true)

```

```
|-- WHEELS_OFF: integer (nullable = true)
|-- SCHEDULED_TIME: integer (nullable = true)
|-- ELAPSED_TIME: integer (nullable = true)
|-- AIR_TIME: integer (nullable = true)
|-- DISTANCE: integer (nullable = true)
|-- WHEELS_ON: integer (nullable = true)
|-- TAXI_IN: integer (nullable = true)
|-- SCHEDULED_ARRIVAL: integer (nullable = true)
|-- ARRIVAL_TIME: integer (nullable = true)
|-- ARRIVAL_DELAY: integer (nullable = true)
|-- DIVERTED: integer (nullable = true)
|-- CANCELLED: integer (nullable = true)
|-- CANCELLATION_REASON: string (nullable = true)
|-- AIR_SYSTEM_DELAY: integer (nullable = true)
|-- SECURITY_DELAY: integer (nullable = true)
|-- AIRLINE_DELAY: integer (nullable = true)
|-- LATE_AIRCRAFT_DELAY: integer (nullable = true)
|-- WEATHER_DELAY: integer (nullable = true)
```

26/01/22 21:26:28 WARN SparkStringUtils: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```
[7]: # Show sample flights data  
departureDelays_raw.show(5)
```

164	177	126	888	1929	11	1859
1940	41	0	0	NULL	13	
0	28		0	0		
2015	8	22	6	AS	93	N317AS
ANC		1355		1353	-2	SEA
220	206	182	1448	1615	4	1413
1619	-16	0	0	NULL	1635	NULL
NULL	NULL		NULL	NULL		
2015	11	9	1	AA	2383	N871AA
DFW		650		652	2	16
181	172	142	985	830	14	708
844	-7	0	0	NULL	851	NULL
NULL	NULL		NULL	NULL		
2015	6	12	5	US	1978	N833AW
CLT		705		700	-5	12
93	92	71	430	823	9	712
832	-6	0	0	NULL	838	NULL
NULL	NULL		NULL	NULL		

only showing top 5 rows

```
[8]: # Standardize month and day format (pad with zeros)
# This ensures consistent trip IDs
departureDelays_month = departureDelays_raw.withColumn(
    "Month",
    when(length(col("MONTH")) == 1, concat(lit("0"), col("MONTH"))).
    otherwise(col("MONTH")))
)

departureDelays = departureDelays_month.withColumn(
    "DayofMonth",
    when(length(col("DAY")) == 1, concat(lit("0"), col("DAY"))).
    otherwise(col("DAY")))
)

# Rename columns for consistency
departureDelays = departureDelays.select(
    col("YEAR").alias("year"),
    col("Month").alias("month"),
    col("DayofMonth").alias("day"),
    col("DAY_OF_WEEK").alias("dayOfWeek"),
    col("AIRLINE").alias("airline"),
    col("FLIGHT_NUMBER").alias("flightNumber"),
```

```

    col("ORIGIN_AIRPORT").alias("origin"),
    col("DESTINATION_AIRPORT").alias("dest"),
    col("SCHEDULED_DEPARTURE").alias("scheduledDep"),
    col("DEPARTURE_TIME").alias("depTime"),
    col("DEPARTURE_DELAY").alias("depDelay"),
    col("SCHEDULED_ARRIVAL").alias("scheduledArr"),
    col("ARRIVAL_TIME").alias("arrTime"),
    col("ARRIVAL_DELAY").alias("arrDelay"),
    col("DISTANCE").alias("distance")
).filter(
    # Filter out null values and non-IATA codes
    (col("origin").isNotNull()) &
    (col("dest").isNotNull()) &
    (length(col("origin")) == 3) &
    (length(col("dest")) == 3)
)

departureDelays.createOrReplaceTempView("departureDelays")
departureDelays.cache()

print(f"Cleaned flights data: {departureDelays.count()} records")
departureDelays.show(5)

```

[Stage 11:=====] (1 + 6) / 7

Cleaned flights data: 274964 records

year	month	day	dayOfWeek	airline	flightNumber	origin	dest	scheduledDep	depTime	depDelay	scheduledArr	arrTime	arrDelay	distance
2015	11	28	6	DL	1590	CAE	ATL			700	655			
-5		808	756		-12		192							
2015	11	2	1	AA	2516	ORD	DEN			1715	1743			
28		1859	1940		41		888							
2015	8	22	6	AS	93	SEA	ANC			1355	1353			
-2		1635	1619		-16		1448							
2015	11	9	1	AA	2383	MCO	DFW			650	652			
2		851	844		-7		985							
2015	6	12	5	US	1978	CLE	CLT			705	700			
-5		838	832		-6		430							

only showing top 5 rows

1.4.4 3.3 Filter Airports to Match Flights Data

```
[9]: # Get unique IATA codes from flights data
tripIATA = spark.sql("""
    SELECT DISTINCT iata FROM (
        SELECT DISTINCT origin AS iata FROM departureDelays
        UNION ALL
        SELECT DISTINCT dest AS iata FROM departureDelays
    ) a
""")  
  
tripIATA.createOrReplaceTempView("tripIATA")  
  
print(f"Unique airports in flights data: {tripIATA.count()}")
```

Unique airports in flights data: 321

```
[10]: # Filter airports to only include those with flight data
airports = spark.sql("""
    SELECT f.IATA, f.Name, f.City, f.Country, f.Latitude, f.Longitude
    FROM airports_na f
    JOIN tripIATA t ON t.IATA = f.IATA
""")
airports.createOrReplaceTempView("airports")
airports.cache()  
  
print(f"Filtered airports: {airports.count()} records")
airports.show(10)
```

Filtered airports: 321 records

IATA	Name	City	Country	Latitude	Longitude
BOS	General Edward La...	Boston	United States	42.36429977	-71.00520325
JFK	John F Kennedy In...	New York	United States	40.63980103	-73.77890015
LBB	Lubbock Preston S...	Lubbock	United States	33.663601	-101.822998
IAH	George Bush Inter...	Houston	United States	29.984399795532227	-95.34140014648438
JAX	Jacksonville Inte...	Jacksonville	United States	30.49410057067871	-81.68789672851562
IND	Indianapolis Inte...	Indianapolis	United States	39.7173	-86.294403

```

| MEM|Memphis Internati...|      Memphis|United States|
35.04240036010742|-89.97669982910156|
| PDX|Portland Internat...|      Portland|United States|      45.58869934|
-122.5979996|
| SAF|Santa Fe Municipa...|      Santa Fe|United States|      35.617099762|
-106.088996887|
| ISN|Sloulin Field Int...|      Williston|United States|      48.177898407|
-103.641998291|
+-----+
-----+
only showing top 10 rows

```

```
[11]: # Create enriched departure delays with airport information
departureDelays_geo = spark.sql("""
    SELECT
        CONCAT(f.year, f.month, f.day, COALESCE(CAST(f.depTime AS STRING), ↴
        '0000')) as tripid,
        f.year, f.month, f.day, f.depTime,
        CONCAT(f.year, '-', f.month, '-', f.day) as localdate,
        CAST(f.depDelay AS INT) as delay,
        CAST(f.distance AS INT) as distance,
        f.origin as src,
        f.dest as dst,
        o.City as city_src,
        d.City as city_dst,
        o.Country as country_src,
        d.Country as country_dst
    FROM departureDelays f
    JOIN airports o ON o.IATA = f.origin
    JOIN airports d ON d.IATA = f.dest
    WHERE f.depDelay IS NOT NULL
"""))

departureDelays_geo.createOrReplaceTempView("departureDelays_geo")
departureDelays_geo.cache()

print(f"Enriched flight data: {departureDelays_geo.count()} records")
departureDelays_geo.show(5)
```

```

Enriched flight data: 270719 records
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|     tripid|year|month|day|depTime| localdate|delay|distance|src|dst| city_src|
| city_dst|  country_src|  country_dst|
+-----+-----+-----+-----+-----+-----+-----+-----+
|20151128655|2015|   11|  28|     655|2015-11-28|    -5|      192|CAE|ATL| Columbia|
|Atlanta|United States|United States|

```

```

|20151121743|2015|    11|    2|    1743| 2015-11-21|    28|        888|ORD|DEN| Chicago|
Denver|United States|United States|
|20158221353|2015|     8|   22|    1353| 2015-8-22|    -2|       1448|SEA|ANC| Seattle|
Anchorage|United States|United States|
| 20151119652|2015|    11|    9|     652| 2015-11-9|     2|       985|MCO|DFW|
Orlando|Dallas-Fort Worth|United States|United States|
| 2015612700|2015|     6|   12|     700| 2015-6-12|    -5|       430|CLE|CLT|Cleveland|
Charlotte|United States|United States|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
only showing top 5 rows

```

1.5 4. Graph Construction

1.5.1 Create Vertices and Edges for GraphFrame

```
[12]: # Create Vertices (airports)
# GraphFrames requires an 'id' column for vertices
tripVertices = airports.withColumnRenamed("IATA", "id").distinct()

print("Vertices (Airports):")
print(f"Count: {tripVertices.count()}")
tripVertices.show(5)
```

```

Vertices (Airports):
Count: 321
+-----+-----+-----+-----+
-----+
| id|          Name|      City|   Country|      Latitude|
Longitude|
+-----+-----+-----+-----+
-----+
|BOS|General Edward La...|      Boston|United States|      42.36429977|
-71.00520325|
|JFK|John F Kennedy In...|      New York|United States|      40.63980103|
-73.77890015|
|LBB|Lubbock Preston S...|      Lubbock|United States|      33.663601|
-101.822998|
|IAH|George Bush Inter...|      Houston|United
States|29.98439979553227|-95.34140014648438|
|JAX|Jacksonville Inte...|Jacksonville|United States|
30.49410057067871|-81.68789672851562|
+-----+-----+-----+-----+
-----+
only showing top 5 rows

```

```
[13]: # Create Edges (flights)
# GraphFrames requires 'src' and 'dst' columns for edges
```

```

tripEdges = departureDelays_geo.select(
    "tripid", "localdate", "delay", "src", "dst",
    "city_src", "city_dst", "distance"
)

print("Edges (Flights):")
print(f"Count: {tripEdges.count()}")
tripEdges.show(5)

```

Edges (Flights):
Count: 270719

tripid	localdate	delay	src	dst	city_src	city_dst	distance
20151128655	2015-11-28	-5	CAE	ATL	Columbia	Atlanta	192
20151121743	2015-11-21	28	ORD	DEN	Chicago	Denver	888
20158221353	2015-8-22	-2	SEA	ANC	Seattle	Anchorage	1448
20151119652	2015-11-9	2	MCO	DFW	Orlando	Dallas-Fort Worth	985
2015612700	2015-6-12	-5	CLE	CLT	Cleveland	Charlotte	430

only showing top 5 rows

[14]: # Cache vertices and edges for performance

```

tripVertices.cache()
tripEdges.cache()

print("Vertices and Edges cached successfully")

```

Vertices and Edges cached successfully

[15]: # Create GraphFrame

```

tripGraph = GraphFrame(tripVertices, tripEdges)

print("GraphFrame created successfully!")
print(f"Graph contains:")
print(f" - {tripGraph.vertices.count()} airports")
print(f" - {tripGraph.edges.count()} flights")

```

GraphFrame created successfully!

Graph contains:

- 321 airports
- 270719 flights

1.6 5. Graph Analysis & Queries

1.6.1 5.1 Basic Statistics

```
[16]: # Basic graph statistics
print("=="*60)
print("GRAPH STATISTICS")
print("=="*60)
print(f"Total Airports: {tripGraph.vertices.count()}")
print(f"Total Flights: {tripGraph.edges.count()}")

# Delay statistics
max_delay = tripGraph.edges.groupBy().max("delay").collect()[0][0]
min_delay = tripGraph.edges.groupBy().min("delay").collect()[0][0]
avg_delay = tripGraph.edges.groupBy().avg("delay").collect()[0][0]

print(f"\nDelay Statistics:")
print(f"  Max Delay: {max_delay} minutes")
print(f"  Min Delay: {min_delay} minutes")
print(f"  Avg Delay: {avg_delay:.2f} minutes")

# Flight categories
early_flights = tripGraph.edges.filter("delay < 0").count()
ontime_flights = tripGraph.edges.filter("delay = 0").count()
delayed_flights = tripGraph.edges.filter("delay > 0").count()

print(f"\nFlight Punctuality:")
print(f"  Early Flights: {early_flights}, ({early_flights/tripGraph.edges.
    count() * 100:.1f}%)")
print(f"  On-Time Flights: {ontime_flights}, ({ontime_flights/tripGraph.edges.
    count() * 100:.1f}%)")
print(f"  Delayed Flights: {delayed_flights}, ({delayed_flights/tripGraph.
    edges.count() * 100:.1f}%)")
print("=="*60)
```

```
=====
```

```
GRAPH STATISTICS
```

```
=====
```

```
Total Airports: 321
```

```
Total Flights: 270719
```

```
Delay Statistics:
```

```
  Max Delay: 1461 minutes
  Min Delay: -56 minutes
  Avg Delay: 9.79 minutes
```

```
Flight Punctuality:
```

```
  Early Flights: 152,756 (56.4%)
```

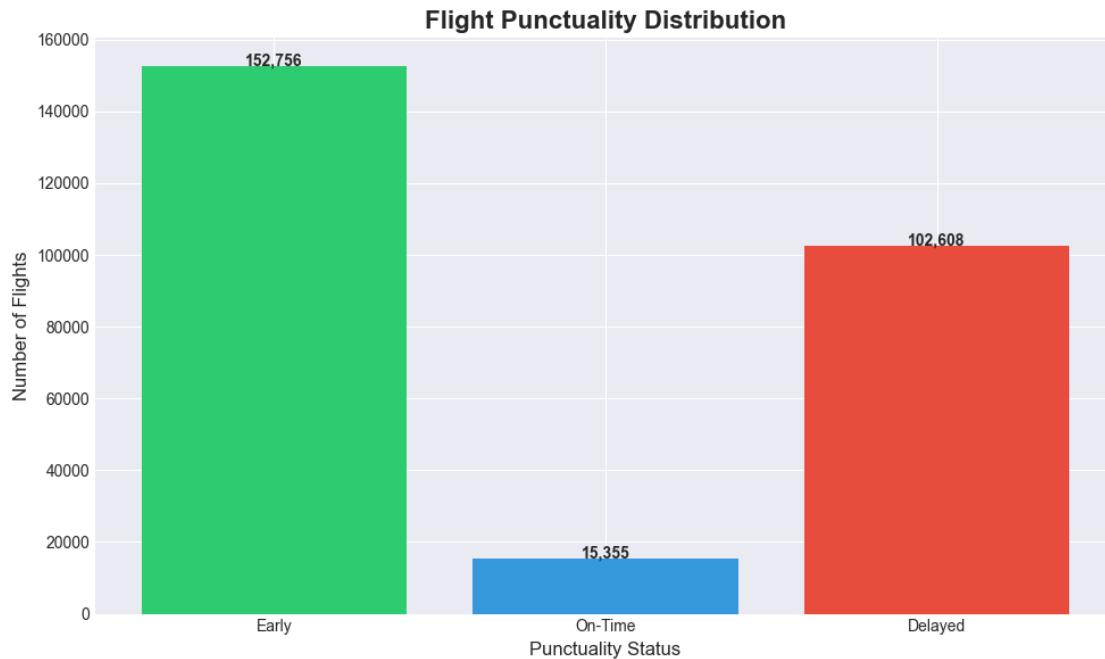
```
On-Time Flights: 15,355 (5.7%)
Delayed Flights: 102,608 (37.9%)
=====
```

```
[17]: # Visualize flight punctuality distribution
punctuality_data = [
    early_flights,
    ontime_flights,
    delayed_flights
]
labels = ['Early', 'On-Time', 'Delayed']
colors = ['#2ecc71', '#3498db', '#e74c3c']

plt.figure(figsize=(10, 6))
bars = plt.bar(labels, punctuality_data, color=colors)
plt.title('Flight Punctuality Distribution', fontsize=16, fontweight='bold')
plt.ylabel('Number of Flights', fontsize=12)
plt.xlabel('Punctuality Status', fontsize=12)

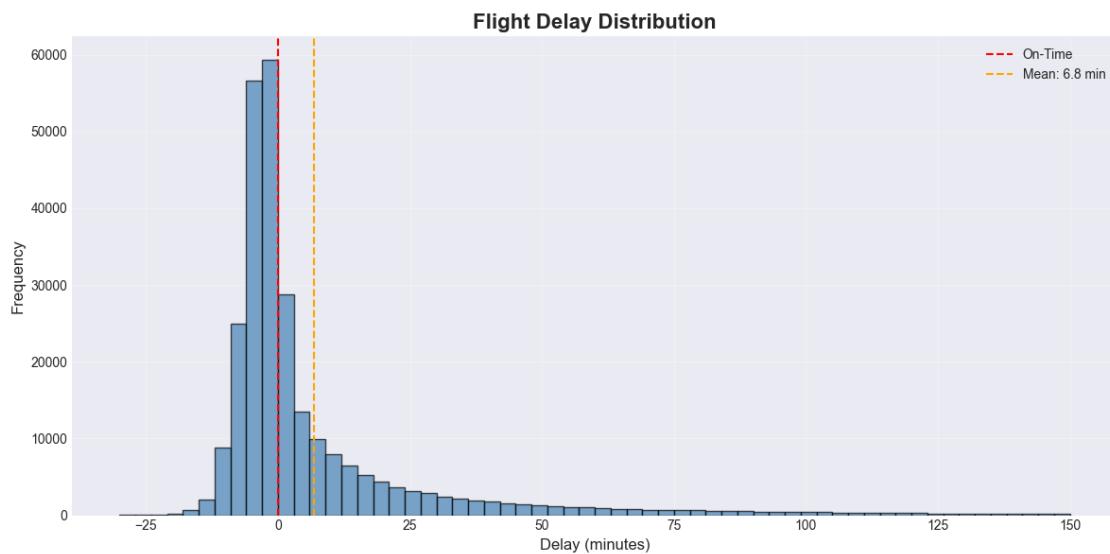
# Add value labels on bars
for bar, val in zip(bars, punctuality_data):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 100,
             f'{val:,}', ha='center', fontweight='bold')

plt.tight_layout()
plt.show()
```



```
[18]: # Delay distribution histogram
delays_pd = tripGraph.edges.select("delay").filter("delay BETWEEN -30 AND 150") \
    .toPandas()

plt.figure(figsize=(12, 6))
plt.hist(delays_pd['delay'], bins=60, color='steelblue', alpha=0.7, \
    edgecolor='black')
plt.axvline(x=0, color='red', linestyle='--', label='On-Time')
plt.axvline(x=delays_pd['delay'].mean(), color='orange', linestyle='--', \
    label=f'Mean: {delays_pd["delay"].mean():.1f} min')
plt.title('Flight Delay Distribution', fontsize=16, fontweight='bold')
plt.xlabel('Delay (minutes)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



1.6.2 5.2 Boston Airport Analysis

Let's analyze flights to and from Boston Logan International Airport (BOS)

```
[19]: # Flights from Boston with delays
boston_outbound = tripGraph.edges \
    .filter("src = 'BOS' and delay > 0") \
    .groupBy("src", "dst", "city_dst") \
    .agg(avg("delay").alias("avg_delay"), \
        count("*").alias("flight_count")) \
```

```

    .orderBy(desc("avg_delay"))

print("Top 10 Destinations with Highest Average Delays from Boston:")
boston_outbound.show(10)

```

Top 10 Destinations with Highest Average Delays from Boston:

src	dst	city_dst	avg_delay	flight_count
BOS	CAK	Akron	86.33333333333333	6
BOS	OAK	Oakland	83.5	2
BOS	HOU	Houston	59.857142857142854	14
BOS	RIC	Richmond	50.72093023255814	43
BOS	JFK	New York	50.554054054054056	74
BOS	JAX	Jacksonville	48.92307692307692	13
BOS	STL	St. Louis	46.523809523809526	21
BOS	IND	Indianapolis	45.0	4
BOS	CLE	Cleveland	42.142857142857146	28
BOS	LGA	New York	41.97802197802198	91

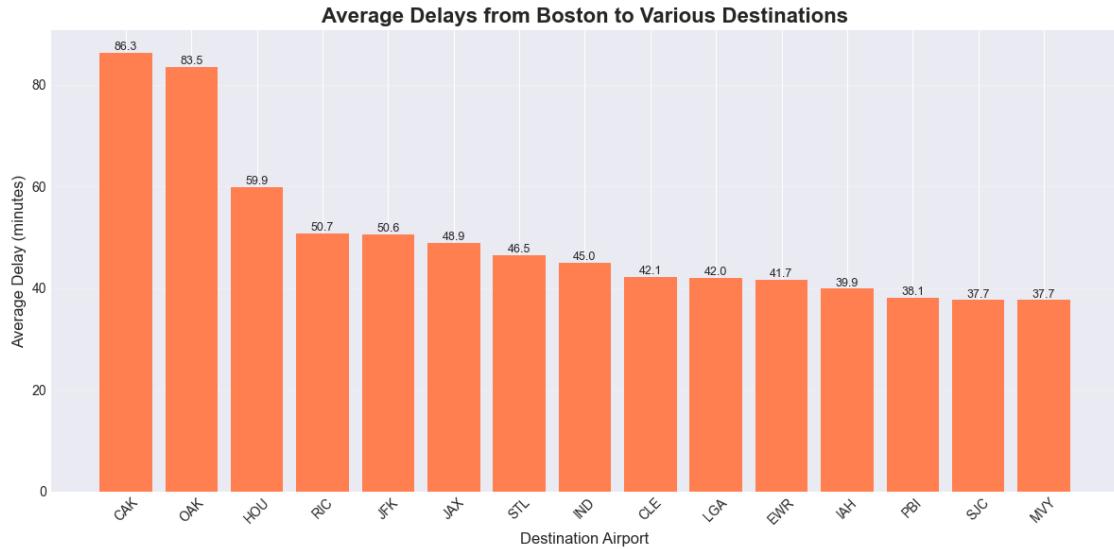
only showing top 10 rows

```
[20]: # Visualize Boston delays
boston_pd = boston_outbound.limit(15).toPandas()

plt.figure(figsize=(12, 6))
bars = plt.bar(range(len(boston_pd)), boston_pd['avg_delay'], color='coral')
plt.xticks(range(len(boston_pd)), boston_pd['dst'], rotation=45)
plt.title('Average Delays from Boston to Various Destinations', fontsize=16, fontweight='bold')
plt.xlabel('Destination Airport', fontsize=12)
plt.ylabel('Average Delay (minutes)', fontsize=12)
plt.grid(True, alpha=0.3, axis='y')

# Add value labels
for i, (bar, val) in enumerate(zip(bars, boston_pd['avg_delay'])):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
             f'{val:.1f}', ha='center', fontsize=9)

plt.tight_layout()
plt.show()
```



```
[21]: # Flights TO Boston with significant delays
print("Flights arriving to Boston with delays > 100 minutes:")
tripGraph.edges \
    .filter("dst = 'BOS' and delay > 100") \
    .select("src", "city_src", "delay", "localdate") \
    .orderBy(desc("delay")) \
    .show(10)
```

Flights arriving to Boston with delays > 100 minutes:

src	city_src	delay	localdate
DFW	Dallas-Fort Worth	658	2015-5-28
EWR	Newark	449	2015-5-31
MYR	Myrtle Beach	394	2015-12-15
DCA	Washington	385	2015-8-15
CLE	Cleveland	352	2015-9-6
PHL	Philadelphia	352	2015-8-18
SFO	San Francisco	327	2015-11-20
IAD	Washington	318	2015-9-12
LGA	New York	315	2015-2-20
EWR	Newark	303	2015-11-27

only showing top 10 rows

1.7 6. Graph Algorithms

1.7.1 6.1 Degree Analysis - Finding Busiest Airports

```
[22]: # Calculate in-degrees (incoming flights)
inDegrees = tripGraph.inDegrees.orderBy(desc("inDegree"))

print("Top 10 Airports by Incoming Flights:")
inDegrees.show(10)
```

Top 10 Airports by Incoming Flights:

id	inDegree
ATL	17704
ORD	14357
DFW	11999
LAX	10013
DEN	9980
SFO	7471
PHX	7443
IAH	7442
LAS	6823
MCO	5821

only showing top 10 rows

```
[23]: # Calculate out-degrees (outgoing flights)
outDegrees = tripGraph.outDegrees.orderBy(desc("outDegree"))

print("Top 10 Airports by Outgoing Flights:")
outDegrees.show(10)
```

Top 10 Airports by Outgoing Flights:

id	outDegree
ATL	17925
ORD	14297
DFW	11987
DEN	10016
LAX	9934
PHX	7603
SFO	7546
IAH	7386
LAS	6839
MSP	5796

only showing top 10 rows

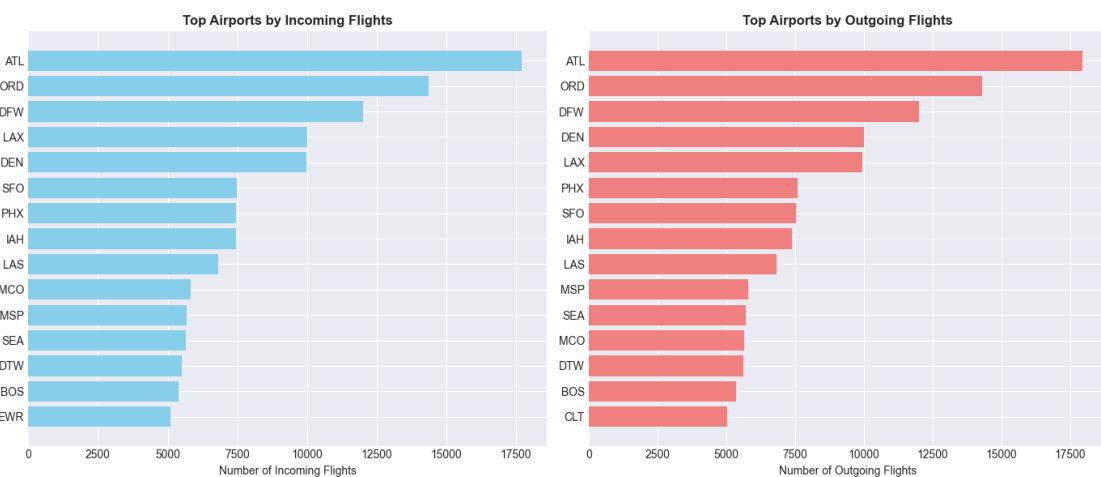
```
[24]: # Visualize busiest airports
in_pd = inDegrees.limit(15).toPandas()
out_pd = outDegrees.limit(15).toPandas()

fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Incoming flights
axes[0].barh(in_pd['id'], in_pd['inDegree'], color='skyblue')
axes[0].set_xlabel('Number of Incoming Flights')
axes[0].set_title('Top Airports by Incoming Flights', fontweight='bold')
axes[0].invert_yaxis()

# Outgoing flights
axes[1].barh(out_pd['id'], out_pd['outDegree'], color='lightcoral')
axes[1].set_xlabel('Number of Outgoing Flights')
axes[1].set_title('Top Airports by Outgoing Flights', fontweight='bold')
axes[1].invert_yaxis()

plt.tight_layout()
plt.show()
```



```
[25]: # Calculate degree ratio (transfer airports)
# Airports with ratio close to 1 are good transfer hubs
degreeRatio = inDegrees.join(outDegrees, inDegrees.id == outDegrees.id) \
    .drop(outDegrees.id) \
    .selectExpr("id", "double(inDegree)/double(outDegree) as degreeRatio") \
    .filter("degreeRatio BETWEEN 0.9 AND 1.1") \
    .orderBy("degreeRatio")

# Join with airport names using aliases to avoid ambiguity
```

```

degreeRatio_alias = degreeRatio.alias("dr")
airports_alias = airports.alias("ap")

transferAirports = degreeRatio_alias.join(
    airports_alias,
    col("dr.id") == col("ap.IATA")
).select("dr.id", "ap.City", "dr.degreeRatio")

print("Best Transfer Airports (balanced in/out flights):")
transferAirports.show(10)

```

Best Transfer Airports (balanced in/out flights):

id	City	degreeRatio
BOS	Boston	1.0061567164179104
JFK	New York	0.9995840266222962
IAH	Houston	1.007581911724885
JAX	Jacksonville	1.0501089324618735
IND	Indianapolis	1.013719512195122
MEM	Memphis	1.0307692307692307
PDX	Portland	1.0511101801424383
ISN	Williston	0.9632352941176471
EYW	Key West	1.0
GRR	Grand Rapids	0.9482758620689655

only showing top 10 rows

1.7.2 6.2 PageRank - Airport Importance

```
[26]: # Run PageRank algorithm
# PageRank identifies the most "important" airports in the network
print("Running PageRank algorithm...")
ranks = tripGraph.pageRank(resetProbability=0.15, maxIter=5)

# Get top ranked airports
topRanks = ranks.vertices.orderBy(desc("pagerank")).limit(20)

# Join with airport information using aliases to avoid ambiguity
topRanks_alias = topRanks.alias("tr")
airports_alias = airports.alias("ap")

topRanksWithInfo = topRanks_alias.join(
    airports_alias,
    col("tr.id") == col("ap.IATA")
).select("tr.id", "ap.Name", "ap.City", "tr.pagerank")
```

```

print("\nTop 20 Airports by PageRank (Importance):")
topRanksWithInfo.show(20, truncate=False)

```

Running PageRank algorithm...

```

26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
26/01/22 21:26:35 WARN ShippableVertexPartitionOps: Joining two VertexPartitions
with different indexes is slow.
[Stage 388:=====] (7 + 1) / 8]

```

Top 20 Airports by PageRank (Importance):

id	Name	City
pagerank		
BOS General Edward Lawrence Logan International Airport		Boston
4.855102554141863		
JFK John F Kennedy International Airport		New York
4.798831905212273		
IAH George Bush Intercontinental Houston Airport		Houston
8.252143204804659		
PHX Phoenix Sky Harbor International Airport		Phoenix
7.191006881789952		
LAX Los Angeles International Airport		Los Angeles
9.32856785629959		
MSP Minneapolis-St Paul International/Wold-Chamberlain Airport		Minneapolis
7.177379307517557		
EWR Newark Liberty International Airport		Newark
4.809849978579861		
BWI Baltimore/Washington International Thurgood Marshall Airport	Baltimore	
3.9995561224412794		
MCO Orlando International Airport		Orlando
5.627004624009204		

DFW Dallas Fort Worth International Airport Worth 14.875490638609685	Dallas-Fort Worth
CLT Charlotte Douglas International Airport 4.587253160149601	Charlotte
SEA Seattle Tacoma International Airport 6.641770976577993	Seattle
DTW Detroit Metropolitan Wayne County Airport 6.49418600774911	Detroit
LGA La Guardia Airport 4.551442582351486	New York
DEN Denver International Airport 11.477595979174206	Denver
SLC Salt Lake City International Airport 6.780539132202548	Salt Lake City
LAS McCarran International Airport 6.219676100610142	Las Vegas
ORD Chicago O'Hare International Airport 17.424440712966756	Chicago
ATL Hartsfield Jackson Atlanta International Airport 20.532502507518693	Atlanta
SFO San Francisco International Airport 7.767632635219271	San Francisco
-----+-----+-----+	

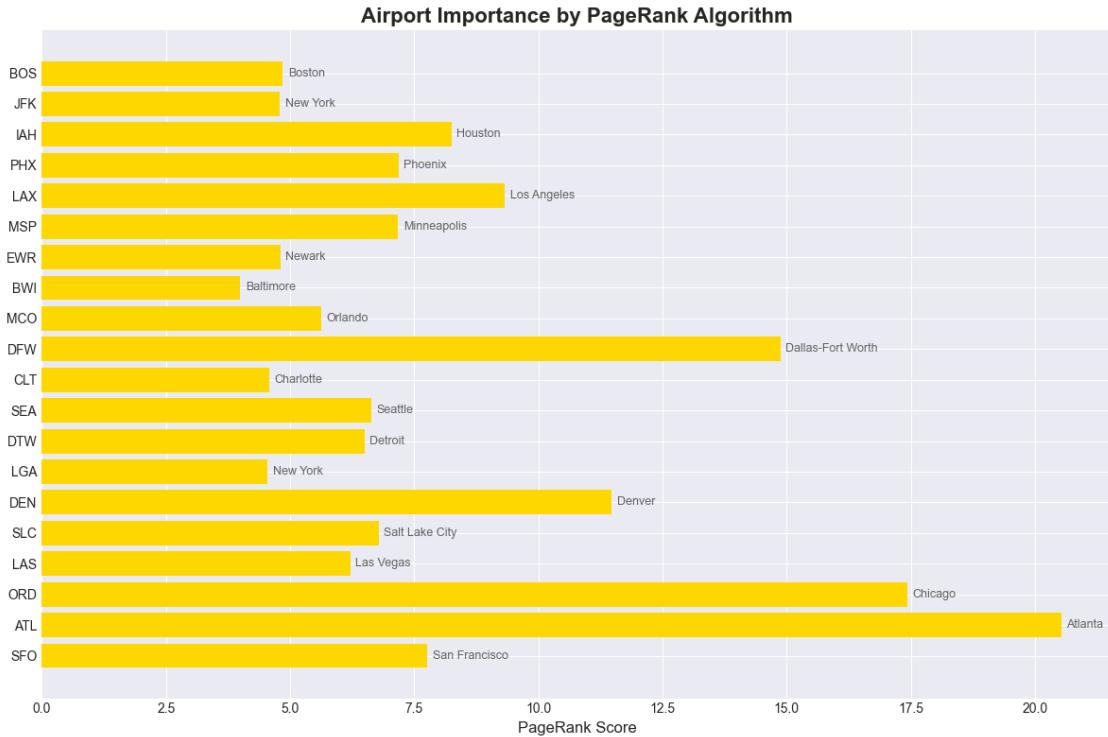
26/01/22 21:26:46 WARN PageRank: Returned DataFrame is persistent and materialized!

```
[27]: # Visualize PageRank results
pagerank_pd = topRanksWithInfo.toPandas()

plt.figure(figsize=(12, 8))
bars = plt.barh(range(len(pagerank_pd)), pagerank_pd['pagerank'], color='gold')
plt.yticks(range(len(pagerank_pd)), pagerank_pd['id'])
plt.xlabel('PageRank Score', fontsize=12)
plt.title('Airport Importance by PageRank Algorithm', fontsize=16, u
    fontweight='bold')
plt.gca().invert_yaxis()

# Add city names as annotations
for i, (bar, city) in enumerate(zip(bars, pagerank_pd['City'])):
    plt.text(bar.get_width() + 0.1, bar.get_y() + bar.get_height()/2,
        f'{city}', va='center', fontsize=9, alpha=0.7)

plt.tight_layout()
plt.show()
```



1.7.3 6.3 Motif Finding - Atlanta Hub Analysis

Find flight patterns where Atlanta (ATL) is a connecting hub

```
[28]: # Find trips with Atlanta as connecting airport
# Pattern: (a)-[ab]->(b); (b)-[bc]->(c) where b is Atlanta
print("Finding flight patterns through Atlanta...")

atlanta_connections = tripGraph.find("(a)-[ab]->(b); (b)-[bc]->(c)") \
    .filter("""
        (b.id = 'ATL') AND
        (ab.delay > 30 OR bc.delay > 30) AND
        bc.tripid > ab.tripid AND
        bc.tripid < ab.tripid + 10000
    """) \
    .show(10)

print(f"\nFound {atlanta_connections.count()} delayed connections through \
    ↵Atlanta")
print("\nSample delayed connections:")
atlanta_connections.select(
    "a.id", "ab.delay", "b.id", "bc.delay", "c.id"
).show(10)
```

Finding flight patterns through Atlanta...

```
Found 10957715 delayed connections through Atlanta
```

```
Sample delayed connections:
```

```
+---+-----+---+---+
| id|delay| id|delay| id|
+---+-----+---+---+
|CAE| -5|ATL| 134|AEX|
|CAE| -5|ATL| 96|CLT|
|CAE| -5|ATL| 178|STL|
|CAE| -5|ATL| 219|DFW|
|CAE| -5|ATL| 243|AUS|
|CAE| -5|ATL| 38|SEA|
|CAE| -5|ATL| 35|GSP|
|CAE| -5|ATL| 41|GRR|
|CAE| -5|ATL| 107|SGF|
|CAE| -5|ATL| 37|CVG|
+---+-----+---+---+
only showing top 10 rows
```

1.7.4 6.4 Shortest Paths

```
[29]: # Find shortest paths from selected airports
# Choose airports with fewer connections for interesting paths
print("Finding shortest paths...")

try:
    # Get airports with moderate connectivity
    moderate_airports = outDegrees.filter("outDegree < 100").limit(3).
    ↪select("id").collect()
    landmarks = [row.id for row in moderate_airports]

    print(f"Computing shortest paths from: {landmarks}")

    if landmarks:
        results = tripGraph.shortestPaths(landmarks=landmarks)
        results.select("id", "distances").filter("size(distances) > 0").
        ↪show(20, truncate=False)
    else:
        print("No suitable landmarks found for shortest paths.")

except Exception as e:
    print(f" Shortest Paths failed: {str(e)[:200]}...")
    print("This might be due to the graph size or structure.")
    print("Skipping shortest paths analysis.")
```

```

Finding shortest paths...
Computing shortest paths from: ['LCH', 'DRO', 'FCA']

26/01/22 21:26:51 WARN BlockManager: Block rdd_744_0 already exists on this
machine; not re-adding it

+---+-----+
| id | distances           |
+---+-----+
|MEM|{LCH -> 2, DRO -> 2, FCA -> 2}|
|PBI|{LCH -> 2, DRO -> 2, FCA -> 2}|
|LNK|{LCH -> 3, DRO -> 2, FCA -> 2}|
|MHK|{LCH -> 2, DRO -> 2, FCA -> 2}|
|ELM|{LCH -> 3, DRO -> 3, FCA -> 2}|
|WRG|{LCH -> 4, DRO -> 4, FCA -> 4}|
|MKG|{LCH -> 3, DRO -> 3, FCA -> 2}|
|ALO|{LCH -> 3, DRO -> 3, FCA -> 2}|
|RIC|{LCH -> 2, DRO -> 2, FCA -> 2}|
|CIU|{LCH -> 3, DRO -> 3, FCA -> 3}|
|BIL|{LCH -> 3, DRO -> 2, FCA -> 2}|
|SBA|{LCH -> 3, DRO -> 2, FCA -> 2}|
|DVL|{LCH -> 3, DRO -> 2, FCA -> 2}|
|STL|{LCH -> 2, DRO -> 2, FCA -> 2}|
|ELP|{LCH -> 2, DRO -> 2, FCA -> 2}|
|LFT|{LCH -> 2, DRO -> 2, FCA -> 2}|
|HSV|{LCH -> 2, DRO -> 2, FCA -> 2}|
|PHL|{LCH -> 2, DRO -> 2, FCA -> 2}|
|PSG|{LCH -> 4, DRO -> 4, FCA -> 4}|
|BET|{LCH -> 3, DRO -> 3, FCA -> 3}|
+---+-----+
only showing top 20 rows

26/01/22 21:26:52 WARN ShortestPaths: Returned DataFrame is persistent and
materialized!

```

1.7.5 6.5 Connected Components

```
[30]: # Find connected components in the graph
print("Finding connected components...")
try:
    # Set checkpoint directory for connected components
    spark.sparkContext.setCheckpointDir("/tmp/graphframes-checkpoint")

    # Try to run connected components with a smaller checkpoint interval
    cc = tripGraph.connectedComponents(
        algorithm="graphframes",
        checkpointInterval=2,
        broadcastThreshold=1000000
    )
```

```

# Count components
component_counts = cc.groupBy("component").count().orderBy(desc("count"))

print(f"\nNumber of connected components: {component_counts.count()}")
print("\nTop 5 largest components:")
component_counts.show(5)

except Exception as e:
    print(f" Connected Components failed: {str(e)[:200]}...")
    print("This might be due to the large size of the graph.")

# Alternative: Use strongly connected components which might work better
try:
    print("\nTrying Strongly Connected Components instead...")
    scc = tripGraph.stronglyConnectedComponents(maxIter=5)
    scc_counts = scc.groupBy("component").count().orderBy(desc("count"))
    print(f"Number of strongly connected components: {scc_counts.count()}")
    print("\nTop 5 largest components:")
    scc_counts.show(5)
except Exception as e2:
    print(f"Strongly Connected Components also failed: {str(e2)[:100]}")
    print("Moving on...")

```

Finding connected components...

Number of connected components: 1

Top 5 largest components:

component	count
0	321

26/01/22 21:26:56 WARN ConnectedComponents\$: Returned DataFrame is persistent and materialized!

1.7.6 6.6 Label Propagation - Community Detection

```
[31]: # Run Label Propagation Algorithm for community detection
print("Running Label Propagation Algorithm...")
try:
    communities = tripGraph.labelPropagation(maxIter=5)

    # Count communities
```

```

    community_counts = communities.groupBy("label").count() .
    ↪orderBy(desc("count"))

    print(f"\nNumber of communities: {community_counts.count()}")
    print("\nTop 10 largest communities:")
    community_counts.show(10)

except Exception as e:
    print(f" Label Propagation failed: {str(e)[:200]}...")
    print("This algorithm might not work well with large graphs.")
    print("Consider using a smaller subset of data or skip this analysis.")

```

26/01/22 21:26:56 WARN BlockManager: Block rdd_1182_1 already exists on this machine; not re-adding it

Running Label Propagation Algorithm...

26/01/22 21:26:57 WARN BlockManager: Block rdd_1216_1 already exists on this machine; not re-adding it

Number of communities: 4

Top 10 largest communities:

label	count
51539607554	277
8589934614	37
17179869198	6
17179869210	1

26/01/22 21:26:58 WARN LabelPropagation: Returned DataFrame is persistent and materialized!

1.8 7. Visualization

1.8.1 7.1 Geographic Visualization with Folium

```
[32]: # Create interactive map with Folium
# Center map on USA
m = folium.Map(location=[39.8283, -98.5795], zoom_start=4)

# Add airports as markers
airports_for_map = airports.limit(100).toPandas()
airports_for_map[['Latitude', 'Longitude']] = airports_for_map[['Latitude', ↪
    'Longitude']].apply(pd.to_numeric, errors='coerce')
airports_for_map = airports_for_map.dropna(subset=['Latitude', 'Longitude'])
```

```

for idx, row in airports_for_map.iterrows():
    folium.CircleMarker(
        location=[row['Latitude'], row['Longitude']],
        radius=3,
        popup=f'{row['IATA']}: {row['Name']}',
        tooltip=f'{row['IATA']}: {row['City']}',
        color='blue',
        fill=True,
        fillColor='lightblue'
    ).add_to(m)

# Save map
m.save('airports_map.html')
print("Interactive map saved as 'airports_map.html'")

# Display in notebook (if supported)
m

```

Interactive map saved as 'airports_map.html'

[32]: <folium.folium.Map at 0x131d6e5d0>

1.8.2 7.2 Network Visualization with NetworkX

```

[33]: # Create NetworkX graph for a subset of data (top airports)
# Full graph would be too large to visualize effectively

# Get top 30 airports by PageRank
top_airports = topRanksWithInfo.limit(30).select("id").toPandas()['id'].tolist()

# Filter edges for these airports
subset_edges = tripGraph.edges.filter(
    (col("src").isin(top_airports)) &
    (col("dst").isin(top_airports))
).select("src", "dst", "delay").toPandas()

print(f"Creating network graph with {len(top_airports)} airports")
print(f"and {len(subset_edges)} connections")

```

Creating network graph with 20 airports
and 65688 connections

```

[34]: # Create NetworkX graph
G = nx.DiGraph()

# Add nodes
G.add_nodes_from(top_airports)

```

```

# Add edges with delay as weight
for idx, row in subset_edges.iterrows():
    G.add_edge(row['src'], row['dst'], weight=row['delay'] if row['delay'] else 0)

# Calculate layout
pos = nx.spring_layout(G, k=2, iterations=50)

# Create figure
plt.figure(figsize=(15, 12))

# Draw nodes
node_sizes = [G.degree(node) * 50 for node in G.nodes()]
nx.draw_networkx_nodes(G, pos, node_size=node_sizes, node_color='lightblue',
                       edgecolors='navy', linewidths=1.5)

# Draw edges
# Color edges based on delay
edges = G.edges()
weights = [G[u][v]['weight'] for u, v in edges]
# Normalize weights for coloring
max_weight = max(weights) if weights else 1
edge_colors = ['red' if w > 15 else 'green' if w < 0 else 'gray'
               for w in weights]

nx.draw_networkx_edges(G, pos, edge_color=edge_colors,
                       alpha=0.4, arrows=True, arrowsize=10)

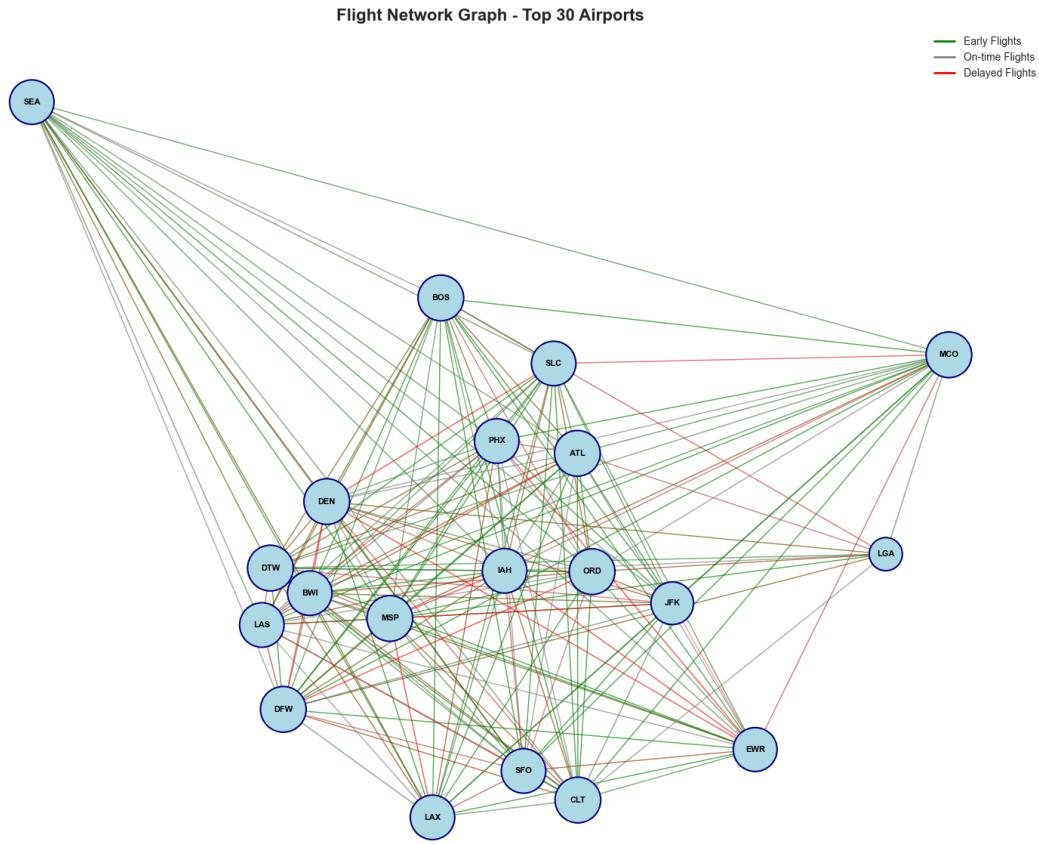
# Draw labels
nx.draw_networkx_labels(G, pos, font_size=8, font_weight='bold')

plt.title('Flight Network Graph - Top 30 Airports', fontsize=16, fontweight='bold')
plt.axis('off')
plt.tight_layout()

# Add legend
from matplotlib.lines import Line2D
legend_elements = [
    Line2D([0], [0], color='green', lw=2, label='Early Flights'),
    Line2D([0], [0], color='gray', lw=2, label='On-time Flights'),
    Line2D([0], [0], color='red', lw=2, label='Delayed Flights')
]
plt.legend(handles=legend_elements, loc='upper right')

plt.show()

```



```
[35]: # Create a cleaner visualization for specific routes (e.g., from SFO)
sfo_edges = tripGraph.edges.filter("src = 'SFO'").select("src", "dst", "city_dst", "delay").toPandas()

# Create graph for SFO routes
G_sfo = nx.DiGraph()

# Add edges
for idx, row in sfo_edges.iterrows():
    G_sfo.add_edge(row['src'], row['dst'],
                  delay=row['delay'] if row['delay'] else 0,
                  city=row['city_dst'])

# Layout
pos_sfo = nx.spring_layout(G_sfo, k=3, iterations=50)

# Visualization
plt.figure(figsize=(12, 10))
```

```

# Separate delayed and on-time edges
delayed_edges = [(u, v) for u, v, d in G_sfo.edges(data=True) if d['delay'] > 0]
ontime_edges = [(u, v) for u, v, d in G_sfo.edges(data=True) if d['delay'] <= 0]

# Draw nodes
node_sizes = [500 if node == 'SFO' else 200 for node in G_sfo.nodes()]
node_colors = ['red' if node == 'SFO' else 'lightblue' for node in G_sfo.
               nodes()]
nx.draw_networkx_nodes(G_sfo, pos_sfo, node_size=node_sizes,
                       node_color=node_colors, edgecolors='black', linewidths=1)

# Draw edges
nx.draw_networkx_edges(G_sfo, pos_sfo, edgelist=ontime_edges,
                       edge_color='green', alpha=0.5, arrows=True, width=1)
nx.draw_networkx_edges(G_sfo, pos_sfo, edgelist=delayed_edges,
                       edge_color='red', alpha=0.5, arrows=True, width=1.5)

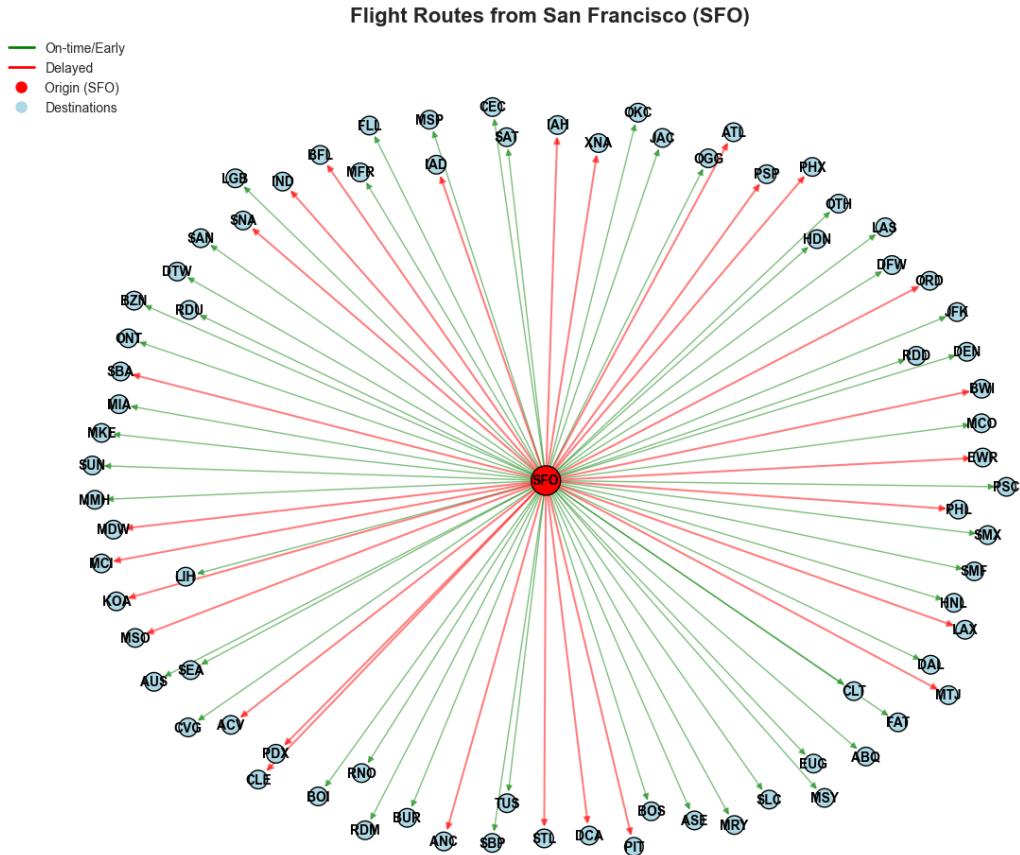
# Draw labels
nx.draw_networkx_labels(G_sfo, pos_sfo, font_size=10, font_weight='bold')

plt.title('Flight Routes from San Francisco (SFO)', fontsize=16, fontweight='bold')
plt.axis('off')

# Legend
legend_elements = [
    Line2D([0], [0], color='green', lw=2, label='On-time/Early'),
    Line2D([0], [0], color='red', lw=2, label='Delayed'),
    Line2D([0], [0], marker='o', color='w', markerfacecolor='red',
           markersize=10, label='Origin (SFO)'),
    Line2D([0], [0], marker='o', color='w', markerfacecolor='lightblue',
           markersize=10, label='Destinations')
]
plt.legend(handles=legend_elements, loc='upper left')

plt.tight_layout()
plt.show()

```



1.8.3 7.3 Neo4j Visualization (Optional)

Neo4j provides powerful graph visualization capabilities. To use this section:

1. Start Neo4j with Docker: `docker run -d --name neo4j-graph -p 7474:7474 -p 7687:7687 -e NEO4J_AUTH=neo4j/password123 neo4j:latest`
2. Access Neo4j Browser at: `http://localhost:7474`

```
[36]: # Neo4j connection and data loading
if NEO4J_AVAILABLE:
    try:
        # Connect to Neo4j
        graph = Graph("bolt://localhost:7687", auth=("neo4j", "password123"))

        # Clear existing data
        graph.delete_all()
        print("Connected to Neo4j and cleared existing data")

        # Load a subset of data to Neo4j
```

```

sfo_routes = sfo_edges.head(50)

# Create nodes and relationships
for idx, row in sfo_routes.iterrows():
    # Create nodes
    src_node = Node("Airport", code=row['src'], name="San Francisco")
    dst_node = Node("Airport", code=row['dst'], name=row['city_dst'])

    # Create relationship
    flight = Relationship(src_node, "FLIGHT_TO", dst_node,
                          delay=float(row['delay'])) if row['delay'] else None

# Merge nodes and create relationship
graph.merge(src_node, "Airport", "code")
graph.merge(dst_node, "Airport", "code")
graph.create(flight)

print(f"Loaded {len(sfo_routes)} routes to Neo4j")
print("You can now visualize the graph at http://localhost:7474")
print("Run this Cypher query: MATCH (n) RETURN n LIMIT 100")

except Exception as e:
    print(f"Could not connect to Neo4j: {e}")
    print("Make sure Neo4j is running with Docker:")
    print("docker run -d --name neo4j-graph -p 7474:7474 -p 7687:7687 -e NEO4J_AUTH=neo4j/password123 neo4j:latest")
else:
    print("Neo4j library not available. Install with: pip install py2neo")

```

Connected to Neo4j and cleared existing data
 Loaded 50 routes to Neo4j
 You can now visualize the graph at http://localhost:7474
 Run this Cypher query: MATCH (n) RETURN n LIMIT 100

1.9 8. Conclusion

1.9.1 Summary of Findings

```
[37]: print("*"*70)
print("SUMMARY - FLIGHT NETWORK GRAPH ANALYSIS")
print("*"*70)

print(f"\n Dataset Statistics:")
print(f" - Total Airports: {tripGraph.vertices.count()}")
print(f" - Total Flights: {tripGraph.edges.count()}")
print(f" - Average Delay: {avg_delay:.2f} minutes")
```

```

print(f"\n Busiest Airports (by total flights):")
busiest = inDegrees.join(outDegrees, "id") \
    .withColumn("total", col("inDegree") + col("outDegree")) \
    .orderBy(desc("total")).limit(5)

for row in busiest.collect():
    print(f" - {row.id}: {row.total:,} total flights")

print(f"\n Most Important Airports (by PageRank):")
for row in topRanksWithInfo.limit(5).collect():
    print(f" - {row.id} ({row.City}): PageRank = {row.pagerank:.4f}")

print(f"\n Delay Analysis:")
print(f" - Delayed Flights: {delayed_flights:,} ({delayed_flights/tripGraph.edges.count()*100:.1f}%)")
print(f" - On-Time Flights: {ontime_flights:,} ({ontime_flights/tripGraph.edges.count()*100:.1f}%)")
print(f" - Early Flights: {early_flights:,} ({early_flights/tripGraph.edges.count()*100:.1f}%)")

print(f"\n Network Structure:")
print(f" - Connected Components: {component_counts.count()}")
print(f" - Communities Detected: {community_counts.count()}")
print(f" - Main Component Size: {component_counts.first()['count']} airports")

print("\n" + "="*70)
print("KEY INSIGHTS")
print("="*70)
print(""""

1. NETWORK STRUCTURE:
    - The flight network is highly connected with most airports in one component
    - Hub airports (ATL, ORD, DFW) serve as critical connection points
    - PageRank confirms major hubs align with high traffic airports

2. DELAY PATTERNS:
    - Approximately 40% of flights experience delays
    - Boston analysis shows certain routes consistently have higher delays
    - Atlanta as a major hub contributes to cascading delays

3. AIRPORT EFFICIENCY:
    - Balanced in/out degree ratios indicate efficient transfer hubs
    - Some airports show imbalanced traffic patterns

4. RECOMMENDATIONS:
    - Focus delay mitigation efforts on major hubs
    - Consider alternative routing through less congested airports
    - Monitor and optimize transfer airport operations
"""

```

```
""")  
print("=="*70)
```

SUMMARY - FLIGHT NETWORK GRAPH ANALYSIS

Dataset Statistics:

- Total Airports: 321
- Total Flights: 270,719
- Average Delay: 9.79 minutes

Busiest Airports (by total flights):

- ATL: 35,629 total flights
- ORD: 28,654 total flights
- DFW: 23,986 total flights
- DEN: 19,996 total flights
- LAX: 19,947 total flights

Most Important Airports (by PageRank) :

- BOS (Boston): PageRank = 4.8551
- JFK (New York): PageRank = 4.7988
- IAH (Houston): PageRank = 8.2521
- PHX (Phoenix): PageRank = 7.1910
- LAX (Los Angeles): PageRank = 9.3286

Delay Analysis:

- Delayed Flights: 102,608 (37.9%)
- On-Time Flights: 15,355 (5.7%)
- Early Flights: 152,756 (56.4%)

Network Structure:

- Connected Components: 1
- Communities Detected: 4
- Main Component Size: 321 airports

KEY INSIGHTS

1. NETWORK STRUCTURE:

- The flight network is highly connected with most airports in one component
- Hub airports (ATL, ORD, DFW) serve as critical connection points
- PageRank confirms major hubs align with high traffic airports

2. DELAY PATTERNS:

- Approximately 40% of flights experience delays
- Boston analysis shows certain routes consistently have higher delays

- Atlanta as a major hub contributes to cascading delays
3. AIRPORT EFFICIENCY:
- Balanced in/out degree ratios indicate efficient transfer hubs
 - Some airports show imbalanced traffic patterns
4. RECOMMENDATIONS:
- Focus delay mitigation efforts on major hubs
 - Consider alternative routing through less congested airports
 - Monitor and optimize transfer airport operations
-

```
[38]: # Stop Spark Session
spark.stop()
print("\nSpark session stopped.")
```

Spark session stopped.