

pyspark_classification

January 22, 2026

1 PySpark Classification dengan Random Forest

Nama: Muhammad Rizky Hajar

NIM: 24.55.2714

Prodi: PJJ Informatika

Konsentrasi: Big Data dan Predictive Analytics (BDPA)

Mata Kuliah: IT Platform and Infrastructure

1.1 Tugas 2: PySpark Hadoop Classification

Konteks Tugas: - Menggunakan PySpark dengan Hadoop - Klasifikasi dengan 4 kelas - Algoritma klasifikasi selain Logistic Regression (menggunakan **Random Forest**) - Dataset: Wine Quality (Red Wine) dari UCI ML Repository

1.2 Wine Quality Dataset - 4 Class Classification

Notebook ini mendemonstrasikan klasifikasi multiclass menggunakan PySpark MLlib dengan algoritma Random Forest.

1.2.1 Struktur:

1. Setup & Import
2. Load & Explore Data
3. Preprocessing
4. Model Training
5. Evaluation
6. Feature Importance

1.3 1. Setup & Import

```
[1]: # Import libraries
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, when, count
from pyspark.sql.types import DoubleType, IntegerType
```

```

# MLlib imports
from pyspark.ml import Pipeline
from pyspark.ml.feature import VectorAssembler, StringIndexer
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

# Visualization
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

print("Libraries imported successfully!")

```

Libraries imported successfully!

```
[2]: # Initialize SparkSession
spark = SparkSession.builder \
    .appName("WineQualityClassification") \
    .config("spark.driver.memory", "4g") \
    .config("spark.sql.shuffle.partitions", "8") \
    .getOrCreate()

# Set log level to reduce noise
spark.sparkContext.setLogLevel("WARN")

print(f"Spark Version: {spark.version}")
print(f"Spark Session: {spark}")
```

WARNING: Using incubator modules: jdk.incubator.vector
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
26/01/22 18:09:55 WARN Utils: Your hostname, Rizkys-MacBook-Air-M4.local,
resolves to a loopback address: 127.0.0.1; using 192.168.1.154 instead (on
interface en0)
26/01/22 18:09:55 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another
address
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).
26/01/22 18:09:56 WARN NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Spark Version: 4.1.1
Spark Session: <pyspark.sql.session.SparkSession object at 0x1228df2f0>

1.4 2. Load & Explore Data

```
[3]: # Load Wine Quality dataset (semicolon separated)
df = spark.read.csv(
    "../data/winequality-red.csv",
    header=True,
    inferSchema=True,
    sep=";"
)

print(f"Dataset loaded: {df.count()} rows, {len(df.columns)} columns")
```

Dataset loaded: 1599 rows, 12 columns

```
[4]: # Show schema
df.printSchema()
```

```
root
 |-- fixed acidity: double (nullable = true)
 |-- volatile acidity: double (nullable = true)
 |-- citric acid: double (nullable = true)
 |-- residual sugar: double (nullable = true)
 |-- chlorides: double (nullable = true)
 |-- free sulfur dioxide: double (nullable = true)
 |-- total sulfur dioxide: double (nullable = true)
 |-- density: double (nullable = true)
 |-- pH: double (nullable = true)
 |-- sulphates: double (nullable = true)
 |-- alcohol: double (nullable = true)
 |-- quality: integer (nullable = true)
```

```
[5]: # Preview data
df.show(10, truncate=False)
```

fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	pH	sulphates	alcohol	quality
7.4	0.7	0.0	1.9	0.076	11.0					
134.0	0.9978	3.51	0.56	9.4	5	1				
7.8	0.88	0.0	2.6	0.098	25.0					
167.0	0.9968	3.2	0.68	9.8	5	1				
7.8	0.76	0.04	2.3	0.092	15.0					
154.0	0.997	3.26	0.65	9.8	5	1				
11.2	0.28	0.56	1.9	0.075	17.0					
160.0	0.998	3.16	0.58	9.8	6	1				

```

| 7.4          | 0.7          | 0.0          | 1.9          | 0.076        | 11.0
| 34.0         |              | 0.9978      | 3.51|0.56    | 9.4          | 5            |
| 7.4          | 0.66         | 0.0          | 1.8          | 0.075        | 13.0
| 40.0         |              | 0.9978      | 3.51|0.56    | 9.4          | 5            |
| 7.9          | 0.6          | 0.06         | 1.6          | 0.069        | 15.0
| 59.0         |              | 0.9964      | 3.3 |0.46    | 9.4          | 5            |
| 7.3          | 0.65         | 0.0          | 1.2          | 0.065        | 15.0
| 21.0         |              | 0.9946      | 3.39|0.47    | 10.0         | 7            |
| 7.8          | 0.58         | 0.02         | 2.0          | 0.073        | 9.0
| 18.0         |              | 0.9968      | 3.36|0.57    | 9.5          | 7            |
| 7.5          | 0.5          | 0.36         | 6.1          | 0.071        | 17.0
| 102.0        |              | 0.9978      | 3.35|0.8     | 10.5         | 5            |
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
only showing top 10 rows

```

```
[6]: # Statistical summary
df.describe().toPandas()
```

26/01/22 18:10:16 WARN SparkStringUtils: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```
[6]: summary      fixed acidity      volatile acidity      citric acid  \
0   count          1599                  1599                  1599
1   mean    8.319637273295838    0.5278205128205131  0.2709756097560964
2   stddev   1.7410963181276948  0.17905970415353525  0.19480113740531824
3   min             4.6                  0.12                  0.0
4   max            15.9                  1.58                  1.0

      residual sugar      chlorides free sulfur dioxide  \
0           1599          1599                  1599
1  2.5388055034396517  0.08746654158849257  15.874921826141339
2  1.40992805950728  0.047065302010090085  10.46015696980971
3           0.9          0.012                  1.0
4           15.5          0.611                  72.0

      total sulfur dioxide      density      pH  \
0           1599          1599                  1599
1  46.46779237023139  0.9967466791744831  3.311113195747343
2  32.89532447829907  0.0018873339538427265  0.15438646490354271
3           6.0          0.99007                2.74
4           289.0          1.00369                4.01

      sulphates      alcohol      quality
0           1599          1599                  1599
1  0.6581488430268921  10.422983114446502  5.6360225140712945
2  0.1695069795901101  1.0656675818473935  0.8075694397347051
```

3	0.33	8.4	3
4	2.0	14.9	8

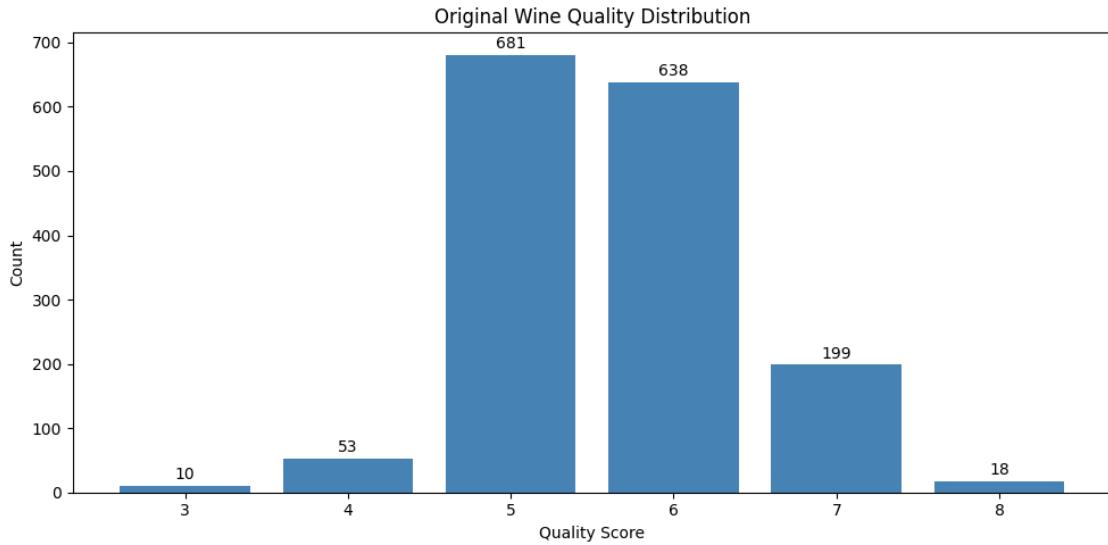
```
[7]: # Check original quality distribution
print("Original Quality Distribution:")
df.groupBy("quality").count().orderBy("quality").show()
```

Original Quality Distribution:

```
+----+----+
|quality|count|
+----+----+
|     3|    10|
|     4|   53|
|     5| 681|
|     6| 638|
|     7| 199|
|     8|   18|
+----+----+
```

```
[8]: # Visualize original quality distribution
quality_counts = df.groupBy("quality").count().orderBy("quality").toPandas()

plt.figure(figsize=(10, 5))
plt.bar(quality_counts['quality'].astype(str), quality_counts['count'], color='steelblue')
plt.xlabel('Quality Score')
plt.ylabel('Count')
plt.title('Original Wine Quality Distribution')
for i, v in enumerate(quality_counts['count']):
    plt.text(i, v + 10, str(v), ha='center')
plt.tight_layout()
plt.show()
```



1.5 3. Preprocessing

1.5.1 3.1 Bucket Quality Score menjadi 4 Kelas

Mapping: - **0 (Bad)**: quality 3-4 - **1 (Average)**: quality 5 - **2 (Good)**: quality 6 - **3 (Excellent)**: quality 7-8

```
[9]: # Create 4-class label from quality score
df_labeled = df.withColumn(
    "label",
    when(col("quality") <= 4, 0) # Bad
    .when(col("quality") == 5, 1) # Average
    .when(col("quality") == 6, 2) # Good
    .otherwise(3) # Excellent (7-8)
)

# Verify new class distribution
print("4-Class Distribution:")
print("0: Bad (3-4), 1: Average (5), 2: Good (6), 3: Excellent (7-8)\n")
df_labeled.groupBy("label").count().orderBy("label").show()
```

4-Class Distribution:

0: Bad (3-4), 1: Average (5), 2: Good (6), 3: Excellent (7-8)

label	count
0	63
1	681
2	638

```
|      3|  217|
+----+----+
```

```
[10]: # Visualize 4-class distribution
class_counts = df_labeled.groupBy("label").count().orderBy("label").toPandas()
class_names = ['Bad (3-4)', 'Average (5)', 'Good (6)', 'Excellent (7-8)']

plt.figure(figsize=(10, 5))
colors = ['#e74c3c', '#f39c12', '#27ae60', '#3498db']
plt.bar(class_names, class_counts['count'], color=colors)
plt.xlabel('Wine Quality Class')
plt.ylabel('Count')
plt.title('4-Class Wine Quality Distribution')
for i, v in enumerate(class_counts['count']):
    plt.text(i, v + 10, str(v), ha='center', fontweight='bold')
plt.tight_layout()
plt.show()
```



1.5.2 3.2 Feature Engineering

```
[11]: # Define feature columns (all columns except quality and label)
feature_columns = [col for col in df.columns if col != 'quality']
print(f"Feature columns ({len(feature_columns)}):")
print(feature_columns)
```

Feature columns (11):
['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH',

```
'sulphates', 'alcohol']

[12]: # Check for null values
print("Null values per column:")
from pyspark.sql.functions import sum as spark_sum, isnan

null_counts = df_labeled.select(
    [spark_sum(col(c).isNull().cast("int")).alias(c) for c in feature_columns]
).toPandas()

print(null_counts.T)
```

```
Null values per column:
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
```

```
[13]: # Create VectorAssembler for features
assembler = VectorAssembler(
    inputCols=feature_columns,
    outputCol="features",
    handleInvalid="skip"
)

# Transform data
df_assembled = assembler.transform(df_labeled)
df_assembled.select("features", "label").show(5, truncate=False)
```

features	label
[7.4,0.7,0.0,1.9,0.076,11.0,34.0,0.9978,3.51,0.56,9.4]	1
[7.8,0.88,0.0,2.6,0.098,25.0,67.0,0.9968,3.2,0.68,9.8]	1
[7.8,0.76,0.04,2.3,0.092,15.0,54.0,0.997,3.26,0.65,9.8]	1
[11.2,0.28,0.56,1.9,0.075,17.0,60.0,0.998,3.16,0.58,9.8]	2
[7.4,0.7,0.0,1.9,0.076,11.0,34.0,0.9978,3.51,0.56,9.4]	1

only showing top 5 rows

1.5.3 3.3 Train/Test Split

```
[14]: # Split data: 80% train, 20% test
train_data, test_data = df_assembled.randomSplit([0.8, 0.2], seed=42)

print(f"Training set: {train_data.count()} samples")
print(f"Test set: {test_data.count()} samples")
```

Training set: 1324 samples
Test set: 275 samples

```
[15]: # Check class distribution in train/test sets
print("Training set class distribution:")
train_data.groupBy("label").count().orderBy("label").show()

print("Test set class distribution:")
test_data.groupBy("label").count().orderBy("label").show()
```

Training set class distribution:

label	count
0	55
1	559
2	532
3	178

Test set class distribution:

label	count
0	8
1	122
2	106
3	39

1.6 4. Model Training - Random Forest

```
[16]: # Initialize Random Forest Classifier
rf = RandomForestClassifier(
    featuresCol="features",
    labelCol="label",
    numTrees=100,           # Number of trees in the forest
    maxDepth=10,            # Maximum depth of each tree
    minInstancesPerNode=5,  # Minimum samples per leaf
```

```

    seed=42
)

print("Random Forest Classifier initialized")
print(f"Parameters:")
print(f" - numTrees: {rf.getNumTrees()}")
print(f" - maxDepth: {rf.getMaxDepth()}")
print(f" - minInstancesPerNode: {rf.getMinInstancesPerNode()}")

```

Random Forest Classifier initialized

Parameters:

- numTrees: 100
- maxDepth: 10
- minInstancesPerNode: 5

```
[17]: # Train the model
print("Training Random Forest model...")
rf_model = rf.fit(train_data)
print("Model training completed!")
```

Training Random Forest model...

```
26/01/22 18:10:47 WARN DAGScheduler: Broadcasting large task binary with size
1570.5 KiB
26/01/22 18:10:47 WARN DAGScheduler: Broadcasting large task binary with size
2.2 MiB
26/01/22 18:10:47 WARN DAGScheduler: Broadcasting large task binary with size
2.9 MiB
26/01/22 18:10:47 WARN DAGScheduler: Broadcasting large task binary with size
3.5 MiB
```

Model training completed!

```
[18]: # Make predictions on test set
predictions = rf_model.transform(test_data)

# Show sample predictions
predictions.select("features", "label", "prediction", "probability").show(10, u
˓→truncate=False)
```

features	label	prediction	probability
[4.9,0.42,0.0,2.1,0.048,16.0,42.0,0.99154,3.71,0.74,14.0]			

```

| 3.0      | [0.0047559088544938035,0.05293779596442267,0.370907717079159,0.57139
85781019245] |
| [5.0,0.74,0.0,1.2,0.041,16.0,46.0,0.99258,4.01,0.59,12.5] | 2
| 2.0      | [0.028304929880715445,0.13531531812365954,0.7441409574352581,0.09223
879456036693] |
| [5.0,1.04,0.24,1.6,0.05,32.0,96.0,0.9934,3.74,0.62,11.5] | 1
| 2.0      | [0.08519678517225249,0.33649681950476146,0.3969867234863248,0.181319
67183666128] |
| [5.2,0.32,0.25,1.8,0.103,13.0,50.0,0.9957,3.38,0.55,9.2] | 1
| 1.0      | [0.04947415961086238,0.5822763712709098,0.35388616455025834,0.014363
304567969358] |
| [5.3,0.47,0.11,2.2,0.048,16.0,89.0,0.99182,3.54,0.88,13.5666666666667] | 3
| 3.0      | [0.007261294261294263,0.06482118777398611,0.10552013709948424,0.8223
973808652354] |
| [5.4,0.42,0.27,2.0,0.092,23.0,55.0,0.99471,3.78,0.64,12.3] | 3
| 2.0      | [0.03702541963723018,0.16643578774921933,0.6990790725771909,0.097459
72003635965] |
| [5.6,0.31,0.37,1.4,0.074,12.0,96.0,0.9954,3.32,0.58,9.2] | 1
| 1.0      | [0.02652256897142962,0.707749771448073,0.25129815852367476,0.0144295
01056822567] |
| [5.6,0.605,0.05,2.4,0.073,19.0,25.0,0.99258,3.56,0.55,12.9] | 1
| 2.0      | [0.05610345173092801,0.28167784785789446,0.5810813309810643,0.081137
36943011336] |
| [5.8,0.29,0.26,1.7,0.063,3.0,11.0,0.9915,3.39,0.54,13.5] | 2
| 2.0      | [0.06833690996915796,0.17369862808268255,0.43053163206011796,0.32743
28298880416] |
| [5.8,0.61,0.11,1.8,0.066,18.0,28.0,0.99483,3.55,0.66,10.9] | 2
| 2.0      |
| [0.02468496029861187,0.3325140658789596,0.4922159106661055,0.15058506315632303]
|
-----+
-----+
-----+
only showing top 10 rows

```

26/01/22 18:10:50 WARN DAGScheduler: Broadcasting large task binary with size
1900.5 KiB

1.7 5. Model Evaluation

```
[19]: # Create evaluators for different metrics
evaluator_accuracy = MulticlassClassificationEvaluator(
    labelCol="label",
    predictionCol="prediction",
    metricName="accuracy"
)

evaluator_f1 = MulticlassClassificationEvaluator(
```

```

        labelCol="label",
        predictionCol="prediction",
        metricName="f1"
    )

evaluator_precision = MulticlassClassificationEvaluator(
    labelCol="label",
    predictionCol="prediction",
    metricName="weightedPrecision"
)

evaluator_recall = MulticlassClassificationEvaluator(
    labelCol="label",
    predictionCol="prediction",
    metricName="weightedRecall"
)

# Calculate metrics
accuracy = evaluator_accuracy.evaluate(predictions)
f1_score = evaluator_f1.evaluate(predictions)
precision = evaluator_precision.evaluate(predictions)
recall = evaluator_recall.evaluate(predictions)

print("=="*50)
print("MODEL EVALUATION RESULTS")
print("=="*50)
print(f"Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f"F1 Score: {f1_score:.4f}")
print(f"Weighted Precision: {precision:.4f}")
print(f"Weighted Recall: {recall:.4f}")
print("=="*50)

```

26/01/22 18:10:52 WARN DAGScheduler: Broadcasting large task binary with size
1910.0 KiB

26/01/22 18:10:52 WARN DAGScheduler: Broadcasting large task binary with size
1910.0 KiB

=====

MODEL EVALUATION RESULTS

=====

Accuracy: 0.6909 (69.09%)

F1 Score: 0.6796

Weighted Precision: 0.6795

Weighted Recall: 0.6909

=====

26/01/22 18:10:53 WARN DAGScheduler: Broadcasting large task binary with size
1910.0 KiB

26/01/22 18:10:53 WARN DAGScheduler: Broadcasting large task binary with size

1910.0 KiB

```
[20]: # Create Confusion Matrix
confusion_df = predictions.groupBy("label", "prediction").count().
    orderBy("label", "prediction").toPandas()

# Pivot to matrix form
confusion_matrix = confusion_df.pivot(index='label', columns='prediction', u
    ↪values='count').fillna(0).astype(int)

print("Confusion Matrix:")
print(confusion_matrix)
```

```
26/01/22 18:10:56 WARN DAGScheduler: Broadcasting large task binary with size
1903.8 KiB
26/01/22 18:10:56 WARN DAGScheduler: Broadcasting large task binary with size
1885.9 KiB
26/01/22 18:10:56 WARN DAGScheduler: Broadcasting large task binary with size
1885.9 KiB
26/01/22 18:10:56 WARN DAGScheduler: Broadcasting large task binary with size
1885.8 KiB
```

Confusion Matrix:

	prediction	1.0	2.0	3.0
label				
0		6	2	0
1		92	29	1
2		23	78	5
3		2	17	20

```
[21]: # Visualize Confusion Matrix
plt.figure(figsize=(10, 8))
class_labels = ['Bad\n(3-4)', 'Average\n(5)', 'Good\n(6)', 'Excellent\n(7-8)']

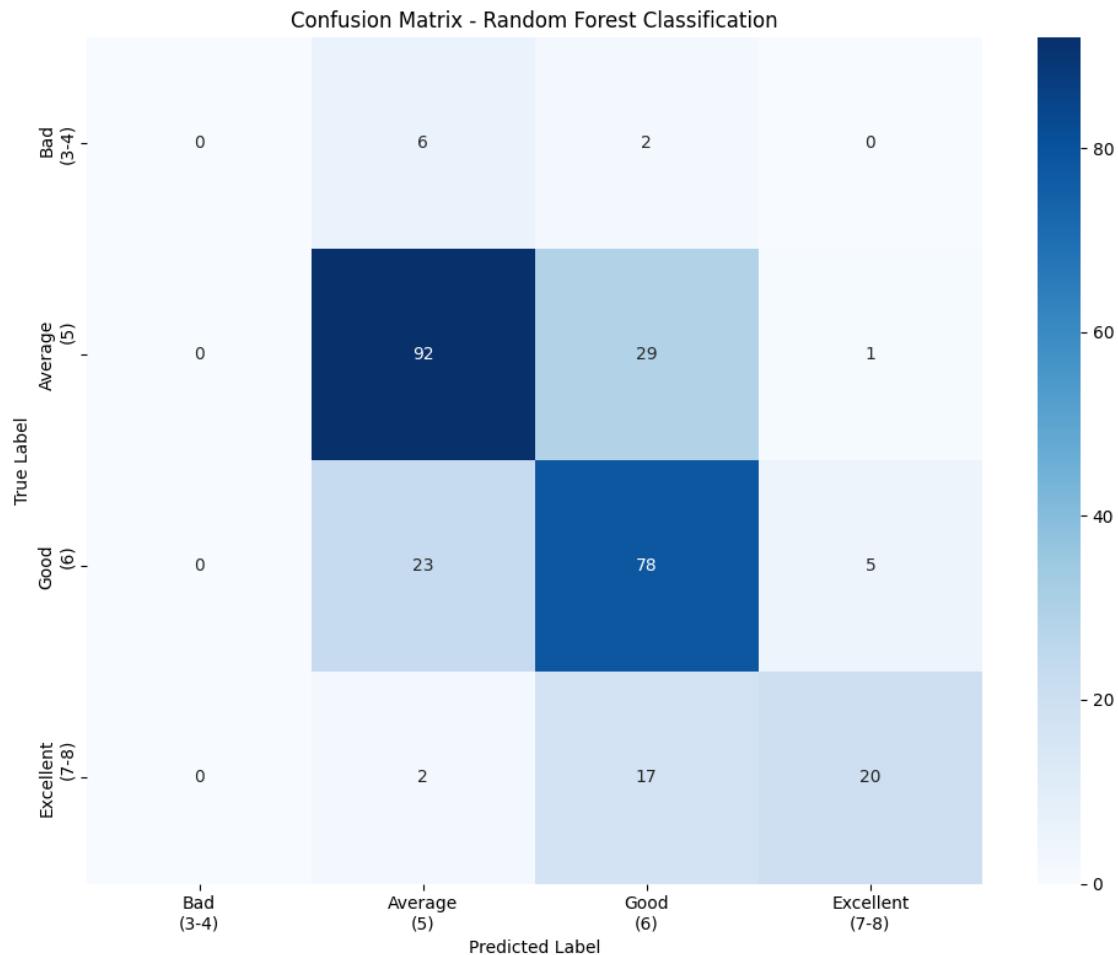
# Ensure all classes are represented
cm_full = np.zeros((4, 4))
for i in range(4):
    for j in range(4):
        if i in confusion_matrix.index and j in confusion_matrix.columns:
            cm_full[i, j] = confusion_matrix.loc[i, j]

sns.heatmap(
    cm_full.astype(int),
    annot=True,
    fmt='d',
    cmap='Blues',
    xticklabels=class_labels,
    yticklabels=class_labels
```

```

)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix - Random Forest Classification')
plt.tight_layout()
plt.show()

```



```

[22]: # Per-class metrics
print("\nPer-Class Performance:")
print("-"*50)

for i, class_name in enumerate(class_labels):
    class_name_clean = class_name.replace('\n', ' ')

```

Calculate per-class metrics

$$\text{tp} = \text{cm_full}[i, i]$$

$$\text{fp} = \text{cm_full}[:, i].\text{sum}() - \text{tp}$$

```

fn = cm_full[i, :].sum() - tp

class_precision = tp / (tp + fp) if (tp + fp) > 0 else 0
class_recall = tp / (tp + fn) if (tp + fn) > 0 else 0
class_f1 = 2 * (class_precision * class_recall) / (class_precision + ↴
class_recall) if (class_precision + class_recall) > 0 else 0

print(f"Class {i} ({class_name_clean}):")
print(f"  Precision: {class_precision:.4f}")
print(f"  Recall:    {class_recall:.4f}")
print(f"  F1-Score:   {class_f1:.4f}")
print()

```

Per-Class Performance:

Class 0 (Bad (3-4)):

```

Precision: 0.0000
Recall:    0.0000
F1-Score:   0.0000

```

Class 1 (Average (5)):

```

Precision: 0.7480
Recall:    0.7541
F1-Score:   0.7510

```

Class 2 (Good (6)):

```

Precision: 0.6190
Recall:    0.7358
F1-Score:   0.6724

```

Class 3 (Excellent (7-8)):

```

Precision: 0.7692
Recall:    0.5128
F1-Score:   0.6154

```

1.8 6. Feature Importance

```
[23]: # Extract feature importance
feature_importance = rf_model.featureImportances.toArray()

# Create DataFrame for visualization
importance_df = pd.DataFrame({
    'feature': feature_columns,
    'importance': feature_importance
}).sort_values('importance', ascending=False)
```

```

print("Feature Importance Ranking:")
print(importance_df.to_string(index=False))

```

Feature Importance Ranking:

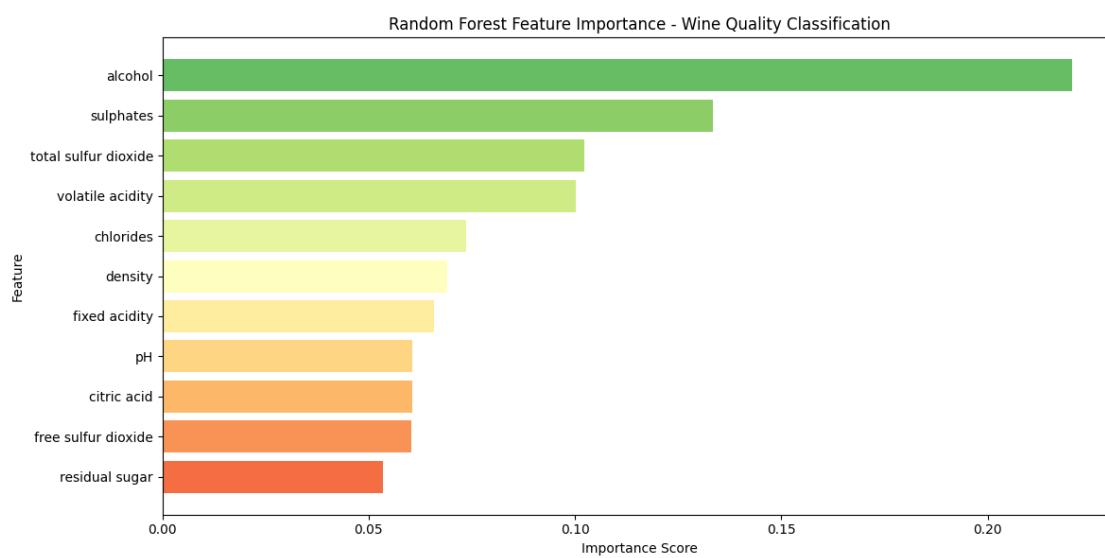
	feature	importance
	alcohol	0.220451
	sulphates	0.133514
total sulfur dioxide		0.102282
	volatile acidity	0.100211
	chlorides	0.073593
	density	0.069050
	fixed acidity	0.065910
	pH	0.060661
	citric acid	0.060479
free sulfur dioxide		0.060315
	residual sugar	0.053534

```

[24]: # Visualize Feature Importance
plt.figure(figsize=(12, 6))

colors = plt.cm.RdYlGn(np.linspace(0.2, 0.8, len(importance_df)))
plt.barh(importance_df['feature'], importance_df['importance'], color=colors[::-1])
plt.xlabel('Importance Score')
plt.ylabel('Feature')
plt.title('Random Forest Feature Importance - Wine Quality Classification')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```



1.9 7. Summary & Conclusions

```
[25]: print("*"*60)
print("SUMMARY - PYSPARK RANDOM FOREST CLASSIFICATION")
print("*"*60)
print(f"\nDataset: Wine Quality (Red Wine)")
print(f"Total Samples: {df.count()}")
print(f"Features: {len(feature_columns)}")
print(f"Classes: 4 (Bad, Average, Good, Excellent)")
print(f"\nTrain/Test Split: 80/20")
print(f"Training Samples: {train_data.count()}")
print(f"Test Samples: {test_data.count()}")
print(f"\nModel: Random Forest Classifier")
print(f" - Number of Trees: 100")
print(f" - Max Depth: 10")
print(f"\nPerformance Metrics:")
print(f" - Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f" - F1 Score: {f1_score:.4f}")
print(f" - Precision: {precision:.4f}")
print(f" - Recall: {recall:.4f}")
print(f"\nTop 3 Important Features:")
for idx, row in importance_df.head(3).iterrows():
    print(f" - {row['feature']}: {row['importance']:.4f}")
print("*"*60)
```

```
=====
SUMMARY - PYSPARK RANDOM FOREST CLASSIFICATION
=====
```

Dataset: Wine Quality (Red Wine)
Total Samples: 1599
Features: 11
Classes: 4 (Bad, Average, Good, Excellent)

Train/Test Split: 80/20
Training Samples: 1324
Test Samples: 275

Model: Random Forest Classifier
- Number of Trees: 100
- Max Depth: 10

Performance Metrics:
- Accuracy: 0.6909 (69.09%)
- F1 Score: 0.6796
- Precision: 0.6795

- Recall: 0.6909

Top 3 Important Features:

- alcohol: 0.2205
 - sulphates: 0.1335
 - total sulfur dioxide: 0.1023
-

[26]: *# Stop Spark Session*

```
spark.stop()  
print("Spark session stopped.")
```

Spark session stopped.